

# Serving dynamic webpages in less than a millisecond

John Fremlin

2008 November 8

# Contents

This talk is about a web-application framework I built.

Introduction

Benchmarking framework overhead

High performance HTTP under Linux

Details of my implementation

Conclusion and final demo

# A web-application framework in Lisp

- Designed for performance
- Still a work in progress
- Built from scratch
- Deals with everything from system calls to user sessions
- Integrated JavaScript generation (AJAX)
- Simple persistent in-memory database
- Made in Common Lisp (no C libraries)

## A web-application framework in Lisp

- Designed for performance
- Still a work in progress
- Built from scratch
- Deals with everything from system calls to user sessions
- Integrated JavaScript generation (AJAX)
- Simple persistent in-memory database
- Made in Common Lisp (no C libraries)

# A web-application framework in Lisp

- Designed for performance
- Still a work in progress
- Built from scratch
- Deals with everything from system calls to user sessions
- Integrated JavaScript generation (AJAX)
- Simple persistent in-memory database
- Made in Common Lisp (no C libraries)

# Motivation

The combination of a fast dynamic webserver with modern webbrowser Javascript implementations is an untapped opportunity for ground-breaking interactive web-applications.

- Many fast static webservers exist: nginx, lighttpd
- But massively slower for dynamic content
- Using interpreted languages
- Horrible caching hacks (fragment caches, varnish, etc.)
- Small AJAX requests are increasingly useful

## Demo: simple message-board

- Post messages to a message board
- Watch them appear immediately without reloading the page
- This is hard to scale (Twitter)

# Implementation

- About 70 lines
- Uses AJAX
- Simple in memory database for messages
- Can be modified dynamically



## Data structures

```
(defrecord message  
  (forum-name :index t)  
  text  
  (author :index t)  
  (time :initform (get-universal-time)))
```

```
(defmyclass (forum (:include simple-channel))  
  name)
```

```
(defvar *fora* (list  
  (make-forum :name "Ubuntu")  
  (make-forum :name "Gentoo")  
  (make-forum :name "Debian"))) 
```

# Website definition

```
(with-site (:page-body-start
  (lambda (title)
    (declare (ignore title))
    '(<div :class "header"
      (<h1
        (<A :href (page-link "/tlug")
          :class "inherit"
          (<span :style (css-attrb :color "red") "TLUG") "_demo" ))
        (output-object-to-ml (webapp-frame))))
      :page-head (lambda (title)
        '(<head
          (<title (output-raw-ml ,title))
          (webapp-default-page-head-contents))))))

(defpage "/tlug" ()
  (webapp "Select_forum"
    (webapp-select-one ""
      *fora*
      :display (lambda (forum) (<span (its name forum)))
      :replace
      (lambda (forum)
        (webapp ()
          (webapp-display forum)))))))
```

## Rendering the data-structures

```
(my-defun forum 'object-to-m1 ()
  (<div :class "forum"
    (<h3 (my name))
    (html-action-form "Post_a_message"
      (text)
      (make-message :forum-name (my name)
                    :text text
                    :author (frame-username (webapp-frame))))
    (my notify)
    (values))

  (<div :class "messages"
    (output-object-to-m1
      (datastore-retrieve-indexed 'message 'forum-name (my name))))
  (output-raw-m1 (call-next-method))))

(my-defun message 'object-to-m1 ()
  (<div :class "message"
    (<p (my text) (<span :class "message-attribution"
                        "_by_" (my author) "_at_" (time-string (my time))))))

(defun time-string (ut)
  (multiple-value-bind
    (second minute hour date month year day daylight-p zone)
    (decode-universal-time ut 0)
    (declare (ignore day daylight-p zone))
    (format nil "~4,'0D-~2,'0D-~2,'0D-~2,'0D:~2,'0D-UTC"
      year month date hour minute second)))
```

# Thoughts

Any questions?

# Benchmarking the framework overhead

How many requests per second can be handled on one core?

- Request a page giving a name
- Reply with `<h1>Hello NAME</h1>` (properly escaping NAME)

Tests the overhead of the framework, excluding the database.

```
schedtool -a 1 -e ab -n 10000 -c100
```

```
http://localhost:3001/?name=TLUG
```

The advantage of my framework is that the complex work to determine the content to display can be done in a fast compiled language.

# Benchmarking the framework overhead

How many requests per second can be handled on one core?

- Request a page giving a name
- Reply with `<h1>Hello NAME</h1>` (properly escaping NAME)

Tests the overhead of the framework, excluding the database.

```
schedtool -a 1 -e ab -n 10000 -c100
```

```
http://localhost:3001/?name=TLUG
```

The advantage of my framework is that the complex work to determine the content to display can be done in a fast compiled language.

## Benchmarking the framework overhead

How many requests per second can be handled on one core?

- Request a page giving a name
- Reply with `<h1>Hello NAME</h1>` (properly escaping NAME)

Tests the overhead of the framework, excluding the database.

```
schedtool -a 1 -e ab -n 10000 -c100
```

```
http://localhost:3001/?name=TLUG
```

The advantage of my framework is that the complex work to determine the content to display can be done in a fast compiled language.

# Ruby/C mongrel web-server

- Bare bones webserver that does not even parse the query string
- Often used for Ruby on Rails
- 1844.88 requests/sec

```
require 'mongrel'
require 'cgi'

class SimpleHandler < Mongrel::Handler
  def process(request, response)
    response.start(200) do |head, out|
      head["Content-Type"] = "text/html"
      name = CGI::escapeHTML(CGI::parse(
        request.params['QUERY_STRING']
      )['name'].first)
      out.write("<h1>Hello #{name}</h1>")
    end
  end
end

h = Mongrel::HttpServer.new("0.0.0.0", "3000")
h.register("/", SimpleHandler.new)
h.run.join
```



# PHP

- Lighttpd with FastCGI
- No code cache
- Logging disabled
- 3174.52 requests/sec

```
<h1>Hello <?= $_REQUEST[ 'name' ] ?></h1>
```

# My implementation

- Automatic escaping
- Plenty of parentheses
- 3806.90 requests/second

```
(defpage "/test" (name)
  (<h1 "Hello_" name))
```

```
(launch-io 'accept-forever
  (make-con-listen :port 3000)
  'http-serve)
```


```
(event-loop)
```

# Varnish

- Cache
- Does not do a dynamic request
- Heavily optimized
- A guiding figure for the maximum possible speed, assuming zero cost for the dynamic request
- 6217.97 requests/second

# Results

**Request throughput on one core (request/s)**



# Thoughts

The framework overhead is not too important, provided it is reasonable. *The advantage of my framework is that the complex work to determine the content to display can be done in a fast compiled language.*

# Thoughts

The framework overhead is not too important, provided it is reasonable. *The advantage of my framework is that the complex work to determine the content to display can be done in a fast compiled language.*

# History of high performance HTTP under Linux

In 1999, Microsoft had the fastest web-server with IIS.

- Mindcraft benchmark (1999 April)
- TUX, a kernelspace webserver by Ingo Molnar at RedHat (2000 July)
- Record breaking SPECWeb99 scores. Twice as fast as IIS

By 2001 May, IIS was again slightly faster than TUX

Performance similar to TUX 2 can now be achieved outside the kernel

- Lighter context-switches
- Less copying for IO
- TCP cork

# History of high performance HTTP under Linux

In 1999, Microsoft had the fastest web-server with IIS.

- Mindcraft benchmark (1999 April)
- TUX, a kernelspace webserver by Ingo Molnar at RedHat (2000 July)
- Record breaking SPECWeb99 scores. Twice as fast as IIS

By 2001 May, IIS was again slightly faster than TUX

Performance similar to TUX 2 can now be achieved outside the kernel

- Lighter context-switches
- Less copying for IO
- TCP cork



# History of high performance HTTP under Linux

In 1999, Microsoft had the fastest web-server with IIS.

- Mindcraft benchmark (1999 April)
- TUX, a kernelspace webserver by Ingo Molnar at RedHat (2000 July)
- Record breaking SPECWeb99 scores. Twice as fast as IIS

By 2001 May, IIS was again slightly faster than TUX  
Performance similar to TUX 2 can now be achieved outside the kernel

- Lighter context-switches
- Less copying for IO
- TCP cork

# History of high performance HTTP under Linux

In 1999, Microsoft had the fastest web-server with IIS.

- Mindcraft benchmark (1999 April)
- TUX, a kernelspace webserver by Ingo Molnar at RedHat (2000 July)
- Record breaking SPECWeb99 scores. Twice as fast as IIS

By 2001 May, IIS was again slightly faster than TUX

Performance similar to TUX 2 can now be achieved outside the kernel

- Lighter context-switches
- Less copying for IO
- TCP cork

# History of high performance HTTP under Linux

In 1999, Microsoft had the fastest web-server with IIS.

- Mindcraft benchmark (1999 April)
- TUX, a kernelspace webserver by Ingo Molnar at RedHat (2000 July)
- Record breaking SPECWeb99 scores. Twice as fast as IIS

By 2001 May, IIS was again slightly faster than TUX  
Performance similar to TUX 2 can now be achieved outside the kernel

- Lighter context-switches
- Less copying for IO
- TCP cork

# Zero-copy IO

- Copying data wastes time
- Simple caches are made useless
- `sendfile(2)` solves this for files from disk
- `writew(2)` helps for dynamic content
- TCP checksum

## Lighter context switches

- Linux always had fast syscalls but the pthreads implementation was very slow
- Native POSIX Thread Library (futexes)

But user-level threading will generally be faster

- Enabled with poll(2)
- Traditionally used by IRC daemons
- But does not scale to large numbers of connexions

## Lighter context switches

- Linux always had fast syscalls but the pthreads implementation was very slow
- Native POSIX Thread Library (futexes)

But user-level threading will generally be faster

- Enabled with poll(2)
- Traditionally used by IRC daemons
- But does not scale to large numbers of connexions

## poll does not scale

```
int poll(struct pollfd *fds, nfds_t nfds, int timeout);

struct pollfd {
    int    fd;           /* file descriptor */
    short  events;       /* requested events */
    short  revents;      /* returned events */
};
```

Poll makes an  $O(n)$  data transfer to the kernel for every wait, where  $n$  is the number of connexions

```
int epoll_wait(int epfd, struct epoll_event *events,
               int maxevents, int timeout);
struct epoll_event {
    uint32_t      events; /* Epoll events */
    epoll_data_t  data;   /* User data variable */
};
int epoll_ctl(int epfd, int op, int fd, struct epoll_event *event);
int epoll_create(int size);
```

## poll does not scale

```
int poll(struct pollfd *fds, nfd_t nfd, int timeout);

struct pollfd {
    int    fd;           /* file descriptor */
    short events;        /* requested events */
    short revents;       /* returned events */
};
```

Poll makes an  $O(n)$  data transfer to the kernel for every wait, where  $n$  is the number of connexions

```
int epoll_wait(int epfd, struct epoll_event *events,
               int maxevents, int timeout);
struct epoll_event {
    uint32_t      events;    /* Epoll events */
    epoll_data_t  data;      /* User data variable */
};
int epoll_ctl(int epfd, int op, int fd, struct epoll_event *event);
int epoll_create(int size);
```



## epoll does scale

- Similar to
  - `/dev/epoll` on Solaris
  - `kqueue` on FreeBSD
- Allows edge-triggering
- Annoyingly cannot be used with disk files
- And AIO cannot be used for network sockets

# Handling many simultaneous connections

- One process per connection: slow
- One OS thread per connection: better
- Multiplexing connections inside one thread: fast
  - select: old
  - poll: better
  - epoll: fastest

## TCP cork

- Avoid sending out partial packets for the HTTP header
- Even if it takes some time to generate the body
- Very important for TUX2

But actually heavily detrimental to performance for me

## TCP cork

- Avoid sending out partial packets for the HTTP header
- Even if it takes some time to generate the body
- Very important for TUX2

But actually heavily detrimental to performance for me

## Engineering decisions

- Implemented entirely in Common Lisp
- No C libraries
- Entirely Linux specific
- Can run at reasonable speeds on SBCL and ClozureCL

## Network module

- Many connexions per thread
- Cannot block in any protocol handler
- Uses code-transformer to generate state-machines from code written in an imperative style
- Key system calls: epoll, read, writev

# HTML generation module

Generates chains of strings to send to `writenv(2)` and offers bonus compile-time typo checking

- Misspelled attributes
- Misplaced tags (for example `<li>` in a `<p>`)
- Misspelled CSS properties

# Parenscrip

A library for writing JavaScript in Lisp. I did not develop this.

- Advanced code generation with Lisp-style macros
- Generates predictable, readable JavaScript
- Easy to debug with in Firebug
- Modified to do more work at compile time
- Very handy, because code can be shared between the server-side and client-side (browser)



# cl-irregsexp

- Regular expression engine
- Unusual syntax
- Fast for some things
- Generates native code

# Conclusion

This project was a huge waste of time!  
And it's not finished.

# Conclusion

This project was a huge waste of time!  
And it's not finished.

## Final demo

The combination of a fast dynamic webserver with modern webbrowser Javascript implementations is an untapped opportunity for ground-breaking interactive web-applications.

Any ideas?