Serving dynamic webpages in less than a millisecond

John Fremlin

2008 November 8

Contents

This talk is about a web-application framework I built.

Introduction

Benchmarking framework overhead

High performance HTTP under Linux

Details of my implementation

Conclusion and final demo

A web-application framework in Lisp

- Designed for performance
- Still a work in progress
- Built from scratch
- Deals with everything from system calls to user sessions
- Integrated JavaScript generation (AJAX)
- Simple persistent in-memory database
- Made in Common Lisp (no C libraries)

Motivation

The combination of a fast dynamic webserver with modern webbrowser Javascript implementations is an untapped opportunity for ground-breaking interactive web-applications.

- Many fast static webservers exist: nginx, lighttpd
- But massively slower for dynamic content
- Using interpreted languages
- Horrible caching hacks (fragment caches, varnish, etc.)
- Small AJAX requests are increasingly useful

Demo: simple message-board

- · Post messages to a message board
- Watch them appear immediately without reloading the page
- This is hard to scale (Twitter)

Implementation

- About 70 lines
- Uses AJAX
- Simple in memory database for messages
- Can be modified dynamically

Data structures

```
(defrecord message
  (forum-name :index t)
  text
  (author :index t)
  (time : initform (get-universal-time)))
(defmyclass (forum (:include simple-channel))
    name)
(defvar *fora* (list
                (make-forum :name "Ubuntu")
                (make-forum :name "Gentoo")
                 (make-forum : name "Debian")))
```

Website definition

```
(with-site (:page-body-start
            (lambda(title)
              (declare (ignore title))
              '(< div : class "header"
                      (< h1)
                        (<A : href (page-link "/tlug")
                            :class "inherit"
                            (<span : style (css-attrib : color "red") "TLUG") "_demo" ))
                      (output-object-to-ml (webapp-frame))))
             :page-head (lambda(title)
                          '( < head
                             (<title (output-raw-ml , title))
                              (webapp-default-page-head-contents))))
  (defpage "/tlug" ()
    (webapp "Select_forum"
      (webapp-select-one
                          *fora*
                          : display (lambda(forum) (<span (its name forum)))
                          : replace
                          (lambda (forum)
                            (webapp ()
                              (webapp-display forum))))))))
```

Rendering the data-structures

```
(mv-defun forum 'object-to-ml ()
 (<div : class "forum"
        (<h3 (mv name))
        (html-action-form "Post_a_message"
            (text)
          (make-message : forum-name (my name)
                         :text text
                         : author (frame-username (webapp-frame)))
          (my notify)
          (values))
        (<div : class "messages"
              (output-object-to-ml
               (datastore-retrieve-indexed 'message 'forum-name (my name))))
        (output-raw-ml (call-next-method))))
(my-defun message 'object-to-ml ()
 (<div : class "message"
        (<p (my text) (<span :class "message-attribution"
                              "_by_" (my author) "_at_" (time-string (my time))))))
(defun time-string (ut)
  (multiple-value-bind
        (second minute hour date month year day daylight-p zone)
      (decode-universal-time ut 0)
     declare (ignore day daylight-p zone))
    (format_nil "~4,'0D-~2,'0D-~2,'0D_~2,'0D:~2,'0D:~2,'0D:~2,'0D_UTC"
            vear month date hour minute second )))
```

Thoughts

Any questions?

Benchmarking the framework overhead

How many requests per second can be handled on one core?

- Request a page giving a name
- Reply with <h1>Hello NAME</h1> (properly escaping NAME)

Tests the overhead of the framework, excluding the database. schedtool -a 1 -e ab -n 10000 -c100 http://localhost:3001/?name=TLUG

The advantage of my framework is that the complex work to determine the content to display can be done in a fast compiled language.

- Bare bones webserver that does not even parse the query string
- Often used for Ruby on Rails
- 1844.88 requests/sec

```
require 'mongrel'
require 'cgi'
class SimpleHandler < Mongrel:: HttpHandler
  def process (request, response)
    response.start(200) do | head, out |
      head ["Content-Type"] = "text/html"
      name = CGI::escapeHTML(CGI::parse(
                         request . params [ 'QUERY_STRING']
                         )['name']. first)
      out.write("<h1>Hello_#{name}</h1>")
    end
  end
end
h = Mongrel :: HttpServer.new("0.0.0.0", "3000")
h.register("/", SimpleHandler.new)
h.run.ioin
```

- Lighttpd with FastCGI
- No code cache
- Logging disabled
- 3174.52 requests/sec

```
<h1>Hello <?= _{REQUEST['name']} ?></h1>
```

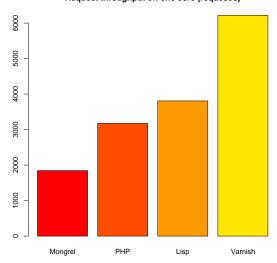
My implementation

- Automatic escaping
- Plenty of parentheses
- 3806.90 requests/second

- Cache
- Does not do a dynamic request
- Heavily optimized
- A guiding figure for the maximum possible speed, assuming zero cost for the dynamic request
- 6217.97 requests/second

Results

Request throughput on one core (request/s)



Thoughts

The framework overhead is not too important, provided it is reasonable. The advantage of my framework is that the complex work to determine the content to display can be done in a fast compiled language.

History of high performance HTTP under Linux

In 1999, Microsoft had the fastest web-server with IIS.

- Mindcraft benchmark (1999 April)
- TUX, a kernelspace webserver by Ingo Molnar at RedHat (2000 July)
- Record breaking SPECWeb99 scores. Twice as fast as IIS

By 2001 May, IIS was again slightly faster than TUX Performance similar to TUX 2 can now be achieved outside the kernel

- Lighter context-switches
- Less copying for IO
- TCP cork

- Copying data wastes time
- Simple caches are made useless
- sendfile(2) solves this for files from disk
- writev(2) helps for dynamic content
- TCP checksum

Lighter context switches

- Linux always had fast syscalls but the pthreads implementation was very slow
- Native POSIX Thread Library (futexes)

But user-level threading will generally be faster

- Enabled with poll(2)
- Traditionally used by IRC daemons
- But does not scale to large numbers of connexions

poll does not scale

```
int poll(struct pollfd *fds, nfds_t nfds, int timeout);
struct pollfd {
  int fd;  /* file descriptor */
short events;  /* requested events */
  short revents; /* returned events */
};
Poll makes an O(n) data transfer to the kernel for every wait,
where n is the number of connexions
int epoll_wait(int epfd, struct epoll_event *events,
  int maxevents, int timeout);
struct epoll_event {
  int epoll_ctl(int epfd, int op, int fd, struct epoll_eve
int epoll_create(int size);
```

- Similar to
 - /dev/epoll on Solaris
 - kqueue on FreeBSD
- Allows edge-triggering
- Annoyingly cannot be used with disk files
- And AIO cannot be used for network sockets

- One process per connection: slow
- One OS thread per connection: better
- Multiplexing connections inside one thread: fast
 - select: oldpoll: betterepoll: fastest

TCP cork

- Avoid sending out partial packets for the HTTP header
- Even if it takes some time to generate the body
- Very important for TUX2

But actually heavily detrimental to performance for me

Engineering decisions

- Implemented entirely in Common Lisp
- No C libraries
- Entirely Linux specific
- Can run at reasonable speeds on SBCL and ClozureCL

Network module

- Many connexions per thread
- Cannot block in any protocol handler
- Uses code-transformer to generate state-machines from code written in an imperative style
- Key system calls: epoll, read, writev

HTML generation module

Generates chains of strings to send to writev(2) and offers bonus compile-time typo checking

- Misspelled attributes
- Misplaced tags (for example in a)
- Misspelled CSS properties

Parenscript

A library for writing JavaScript in Lisp. I did not develop this.

- Advanced code generation with Lisp-style macros
- Generates predictable, readable JavaScript
- Easy to debug with in Firebug
- Modified to do more work at compile time
- Very handy, because code can be shared between the server-side and client-side (browser)

- Regular expression engine
- Unusual syntax
- Fast for some things
- Generates native code

Conclusion

This project was a huge waste of time! And it's not finished.

Final demo

The combination of a fast dynamic webserver with modern webbrowser Javascript implementations is an untapped opportunity for ground-breaking interactive web-applications.

Any ideas?